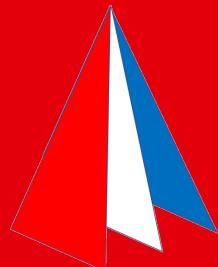


Quick Clean for the Web

Dante van Gemert

Thursday 5th February 2026



Radboud Universiteit



Web is where it's at

E-mail, Google Docs, Typst, custom keyboard software, ...

But: JavaScript 😬

Clean

Functional programming language

Clean

Functional programming language

iTasks

- Structure code with tasks
- Automagically generates web UI ✨

Clean

Functional programming language

iTasks

- Structure code with tasks
- Automagically generates web UI ✨
- Program runs on server
- Small bits of Clean code in browser
for creating UI elements

Clean

Functional programming language

iTasks

- Structure code with tasks
- Automagically generates web UI ✨
- Program runs on server
- Small bits of Clean code in browser
for creating UI elements
- Beneficial to run more in browser
 - Reduce latency, increase availability

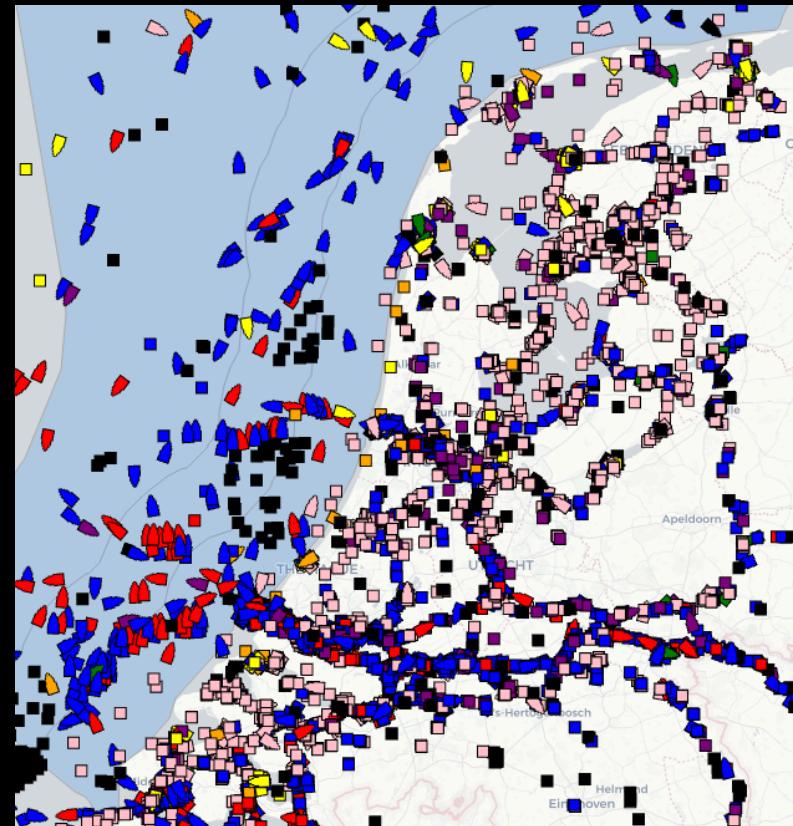
Clean

Functional programming language

iTasks

- Structure code with tasks
- Automagically generates web UI ⚡
- Program runs on server
- Small bits of Clean code in browser for creating UI elements
- Beneficial to run more in browser
 - Reduce latency, increase availability

But: interpreting is too slow



VIIA

Compile instead of interpret

Clean-LLVM

Clean → ?

Compile instead of interpret

In short

- Clean web-app → iTasks
- Better performance → Clean-LLVM (instead of ABC interpreter)

Clean-LLVM

Clean → ?

Compile instead of interpret

In short

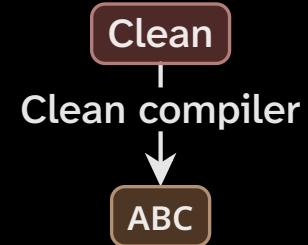
- Clean web-app → iTasks
- Better performance → Clean-LLVM (instead of ABC interpreter)

But: don't want to compile Clean code directly

ABC

Clean → ABC → ?

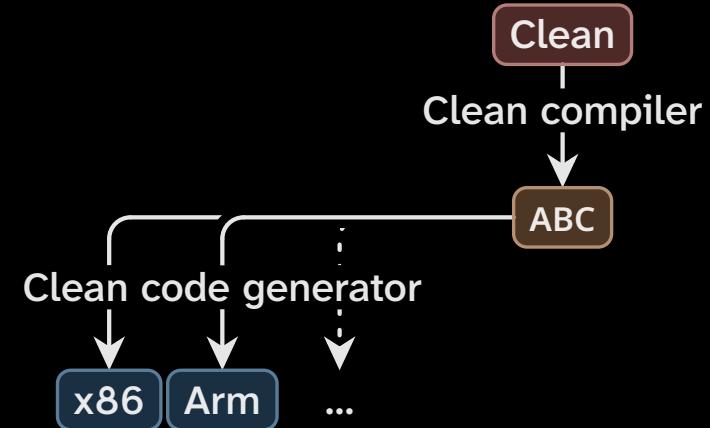
Use Clean compiler's ABC bytecode



ABC

Clean → ABC → ?

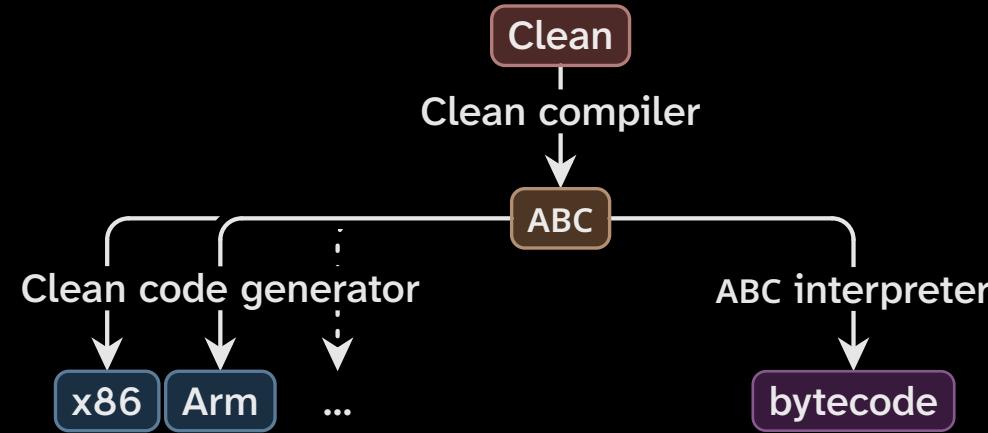
Use Clean compiler's ABC bytecode



ABC

Clean → ABC → ?

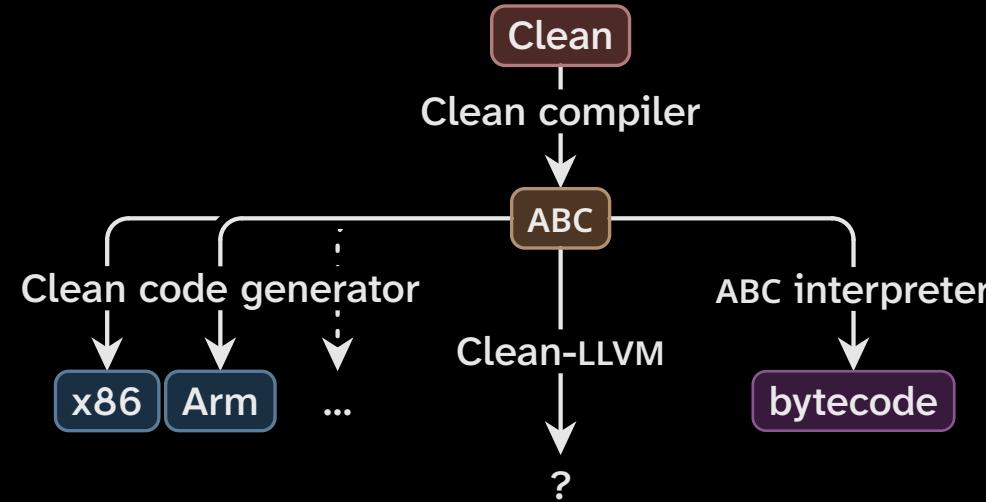
Use Clean compiler's ABC bytecode



ABC

Clean → ABC → ?

Use Clean compiler's ABC bytecode

What do we compile **to?**

WebAssembly (Wasm)

Clean → ABC → ? → WebAssembly

- Compilation target for the web
- Stack based
- Safe & secure (→ restrictive)
- C(++), Rust, Go, Kotlin, Dart, ...

WebAssembly (Wasm)

Clean → ABC → ? → WebAssembly

- Compilation target for the web
- Stack based
- Safe & secure (→ restrictive)
- C(++), Rust, Go, Kotlin, Dart, ...

Originally focussed on imperative languages

- Recently: tail calls

WebAssembly (Wasm)

Clean → ABC → ? → WebAssembly

- Compilation target for the web
- Stack based
- Safe & secure (→ restrictive)
- C(++), Rust, Go, Kotlin, Dart, ...

Originally focussed on imperative languages

- Recently: tail calls

But: we can do better than generating Wasm directly

- Compiler toolkit
- Clang compiler for C & C++
- Intermediate representation (IR)
- Optimisation passes (inlining, dead code elimination, ...)
- Rust, Zig, Julia, Swift, ...

- Compiler toolkit
- Clang compiler for C & C++
- Intermediate representation (IR)
- Optimisation passes (inlining, dead code elimination, ...)
- Rust, Zig, Julia, Swift, ...

Why?

- Generates Wasm, but also x86, Arm, RISC-V...
- Optimisations

- Compiler toolkit
- Clang compiler for C & C++
- Intermediate representation (IR)
- Optimisation passes (inlining, dead code elimination, ...)
- Rust, Zig, Julia, Swift, ...

Why?

- Generates Wasm, but also x86, Arm, RISC-V...
- Optimisations

Wasm is not a standard target for LLVM

- Requires some creative solutions

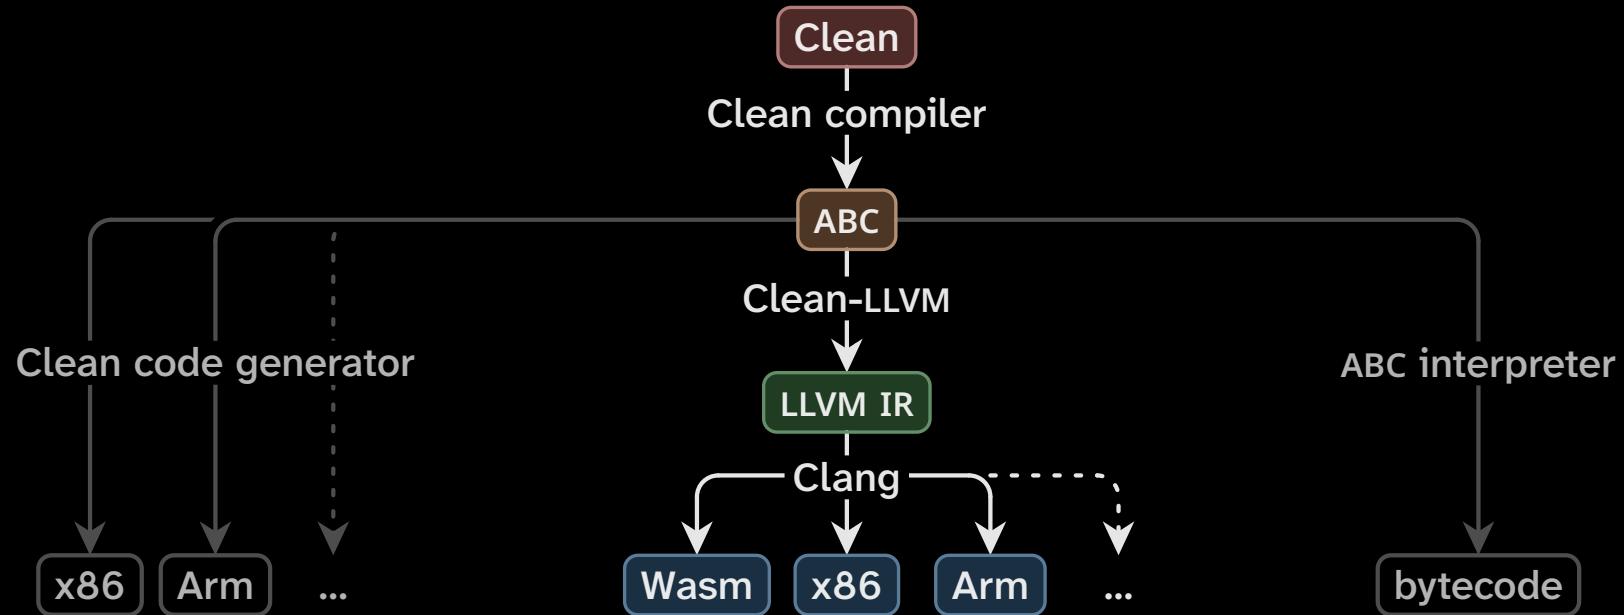
Overview



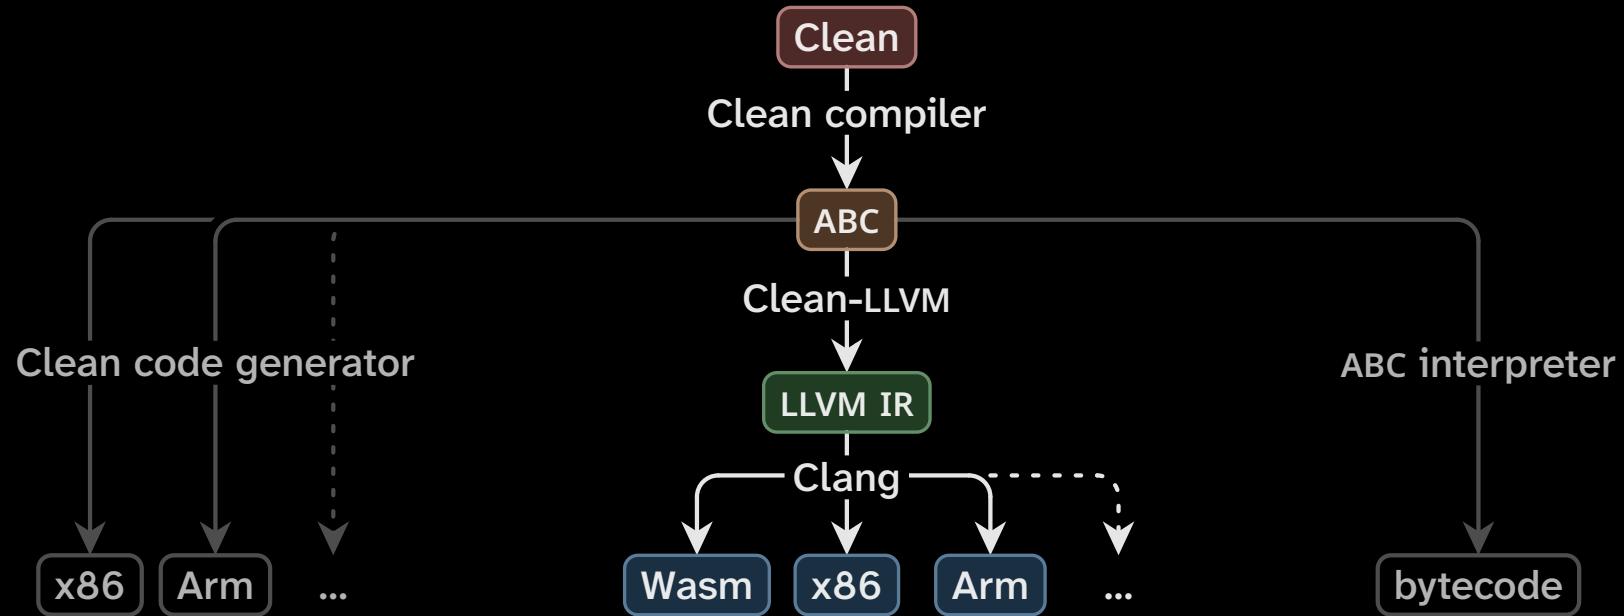
Overview



Overview



Overview



Focus on WebAssembly

GC Integration

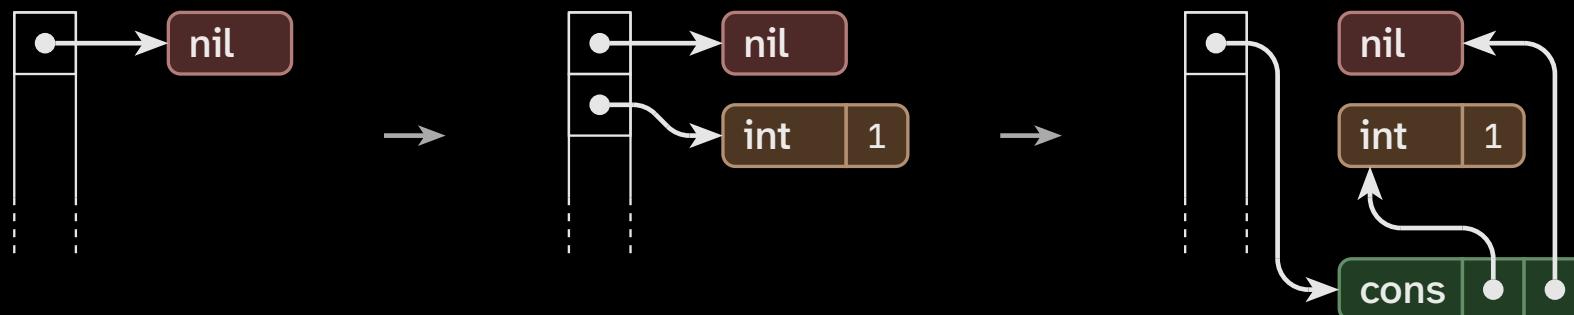
Heap and A Stack

A stack: references to the heap (boxed values)

B stack: basic values (unboxed values: int, char, bool, ...)

C stack: control flow

(hence the name of ABC code)



Creation of the list [1].

GC Integration

- A stack is on WebAssembly stack, for performance
- Can only access stack frame of current function

GC Integration

- A stack is on WebAssembly stack, for performance
- Can only access stack frame of current function

Shadow stack

1. **Spill** (write A stack to memory)
2. **Restore** (read back new addresses)

GC Integration

- A stack is on WebAssembly stack, for performance
- Can only access stack frame of current function

Shadow stack

1. **Spill** (write A stack to memory)
2. **Restore** (read back new addresses)

LLVM passes

- Run optimisations before generating spills/restores: inlining, dead code elimination
- Leads to less function calls → less spill locations

Shadow Stack Overview

mark → statepoints → asp arg → spill

1. Mark A stack (preparation)
2. Insert statepoints
3. Add shadow stack argument
4. Add spills

Preparation: Mark A Stack

mark → statepoints → asp arg → spill

- Which LLVM variables are on A stack?
- Mark them with special type

Preparation: Mark A Stack

mark → statepoints → asp arg → spill

- Which LLVM variables are on A stack?
 - Mark them with special type
1. Mark heap references (A stack values)

ptr %a → *heapptr %a*

Note: simplified / imaginary syntax

LLVM Pass: Insert Statepoints

mark → statepoints → asp arg → spill

- Statepoint: point at which GC *can* run
- Only need to update spilled A stack across statepoints
 - spill before, restore after

LLVM Pass: Insert Statepoints

mark → statepoints → asp arg → spill

- Statepoint: point at which GC *can* run
- Only need to update spilled A stack across statepoints
 - spill before, restore after

2. Convert function calls to statepoints

```
%answer = call i64 @f(42)
```

LLVM Pass: Insert Statepoints

mark → statepoints → asp arg → spill

- Statepoint: point at which GC *can* run
- Only need to update spilled A stack across statepoints
 - spill before, restore after

2. Convert function calls to statepoints

```
%answer = call i64 @f(42)
```



```
%s = call token @gc.statepoint(@f, 42) [ "gc-live"(heapptr %a) ]
```

LLVM Pass: Insert Statepoints

mark → statepoints → asp arg → spill

- Statepoint: point at which GC *can* run
- Only need to update spilled A stack across statepoints
 - spill before, restore after

2. Convert function calls to statepoints

```
%answer = call i64 @f(42)
```



```
%s = call token @gc.statepoint(@f, 42) [ "gc-live"(heapptr %a) ]
```

```
%answer = call i64 @gc.result(%s)
```

LLVM Pass: Insert Statepoints

mark → statepoints → asp arg → spill

- Statepoint: point at which GC *can* run
- Only need to update spilled A stack across statepoints
 - spill before, restore after

2. Convert function calls to statepoints

```
%answer = call i64 @f(42)
```



```
%s = call token @gc.statepoint(@f, 42) [ "gc-live"(heapptr %a) ]
```

```
%answer = call i64 @gc.result(%s)
```

```
%a.new = call i64 @gc.relocate(%s, 0)
```

LLVM Pass: Add Stack Top Argument

mark → statepoints → asp arg → spill

3. Add argument for pointer to spilled A stack

- asp = **A Stack Pointer**

```
define i64 @f(i64 %x)
```



```
define { ptr, i64 } @f_asp(ptr %asp, i64 %x)
```

LLVM Pass: Add Spills

mark → statepoints → asp arg → spill

4. Transform statepoints to function calls, spills, restores

```
%s = call token @gc.statepoint(0f, 42) [ "gc-live"(heapptr %a) ]  
%answer = call i64 @gc.result(%s)
```

LLVM Pass: Add Spills

mark → statepoints → asp arg → spill

4. Transform statepoints to function calls, spills, restores

```
%s = call token @gc.statepoint(@f, 42) [ "gc-live"(heapptr %a) ]  
%answer = call i64 @gc.result(%s)
```

↓

```
store %a to %asp ; spill  
{ %asp.new, %answer } = call i64 @f_asp(%asp, 42)
```

LLVM Pass: Add Spills

mark → statepoints → asp arg → spill

4. Transform statepoints to function calls, spills, restores

```
%a.new = call i64 @gc.relocate(%s, 0)
```



```
%a.new = load %asp.new ; restore
```

LLVM Passes: All Together

mark → statepoints → asp arg → spill

```
%answer = call i64 @f(42)
```

↓

```
store %a to %asp ; spill
```

```
{ %asp.new, %answer } = call i64 @f_asp(%asp, 42)
```

```
%a.new = load %asp.new ; restore
```

LLVM Passes: All Together

mark → statepoints → asp arg → spill

```
%answer = call i64 @f(42)
```

↓

```
store %a to %asp ; spill
```

```
{ %asp.new, %answer } = call i64 @f_asp(%asp, 42)
```

```
%a.new = load %asp.new ; restore
```

Why statepoints?

Implementation for native x86 without shadow stack (theoretically)

Results

Benchmarks

Restriction: not all ABC instructions implemented yet

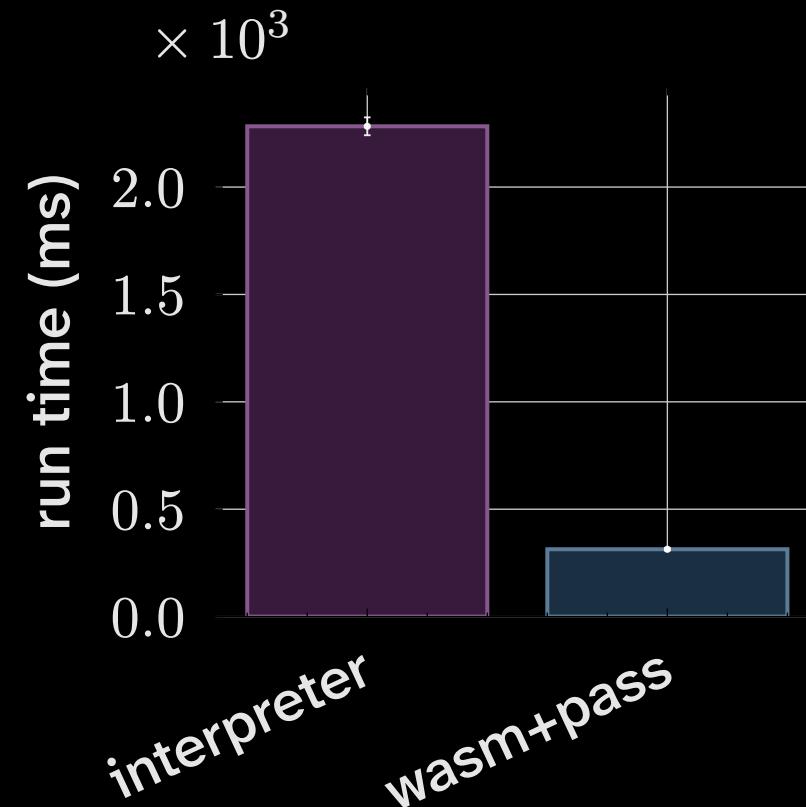
Benchmarks

Restriction: not all ABC instructions implemented yet

- List
 - Create a long list and sum it
- Astack
 - Create 16 heap nodes and pass them to multiple function calls
- **Binarytrees**
 - From the Computer Language Benchmark Game
 - Required additional ABC instructions

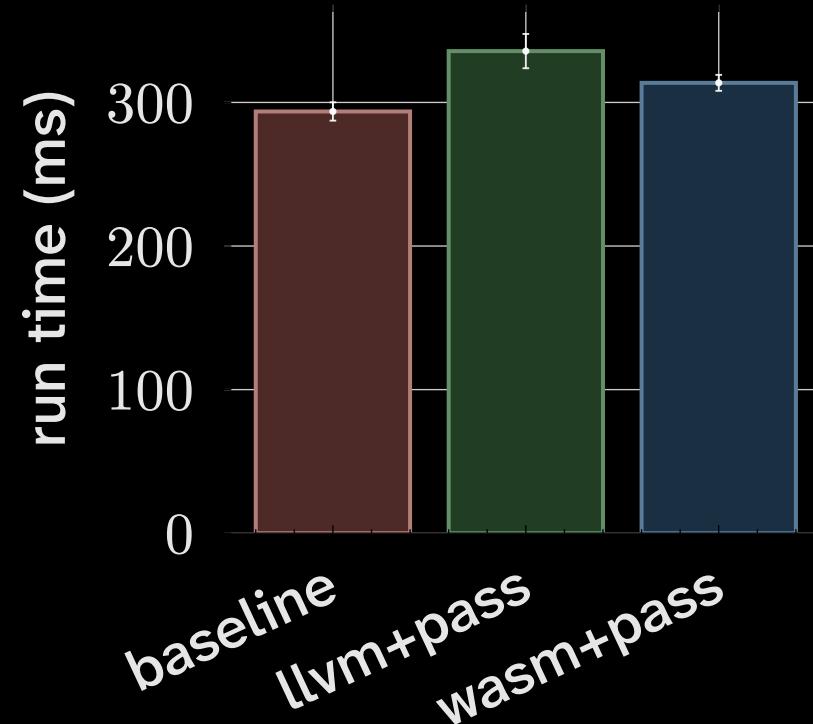
Benchmark: Binarytrees

Wasm: 7 times faster than ABC interpreter (2.3 s vs 0.3 s)



Benchmark: Binarytrees

Close to original Clean compiler's performance



Conclusion

- **Why?** More client-side iTasks code → interpreter too slow

Conclusion

- **Why?** More client-side iTasks code → interpreter too slow
- Clean → ABC → LLVM → WebAssembly

Conclusion

- **Why?** More client-side iTasks code → interpreter too slow
- Clean → ABC → LLVM → WebAssembly
- **GC integration** with shadow A stack: spill/restore

Conclusion

- **Why?** More client-side iTasks code → interpreter too slow
- Clean → ABC → LLVM → WebAssembly
- **GC integration** with shadow A stack: spill/restore
- **LLVM passes** with statepoints
 1. Mark A stack (not a pass, but preparation)
 2. Convert function calls to statepoints
 3. Add shadow A stack pointer argument
 4. Convert statepoints to spills/restores

Conclusion

- **Why?** More client-side iTasks code → interpreter too slow
- Clean → ABC → LLVM → WebAssembly
- **GC integration** with shadow A stack: spill/restore
- **LLVM passes** with statepoints
 1. Mark A stack (not a pass, but preparation)
 2. Convert function calls to statepoints
 3. Add shadow A stack pointer argument
 4. Convert statepoints to spills/restores
- **Optimisation:** inlining before generating statepoints

Conclusion

- **Why?** More client-side iTasks code → interpreter too slow
- Clean → ABC → LLVM → WebAssembly
- **GC integration** with shadow A stack: spill/restore
- **LLVM passes** with statepoints
 1. Mark A stack (not a pass, but preparation)
 2. Convert function calls to statepoints
 3. Add shadow A stack pointer argument
 4. Convert statepoints to spills/restores
- **Optimisation:** inlining before generating statepoints
- **Results:** 7 times faster than interpreter, close to original Clean

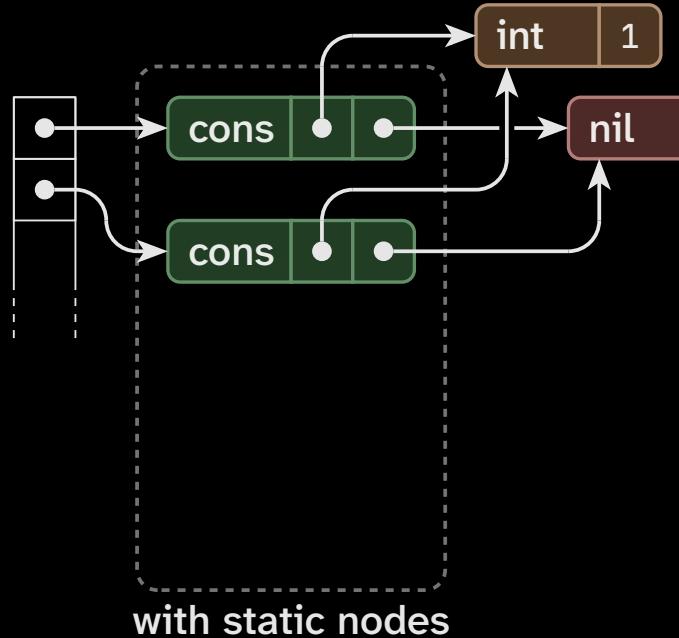
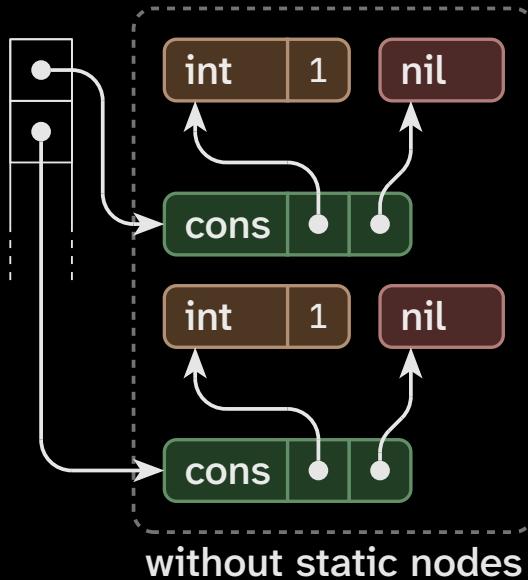
Conclusion

- **Why?** More client-side iTasks code → interpreter too slow
- Clean → ABC → LLVM → WebAssembly
- **GC integration** with shadow A stack: spill/restore
- **LLVM passes** with statepoints
 1. Mark A stack (not a pass, but preparation)
 2. Convert function calls to statepoints
 3. Add shadow A stack pointer argument
 4. Convert statepoints to spills/restores
- **Optimisation:** inlining before generating statepoints
- **Results:** 7 times faster than interpreter, close to original Clean

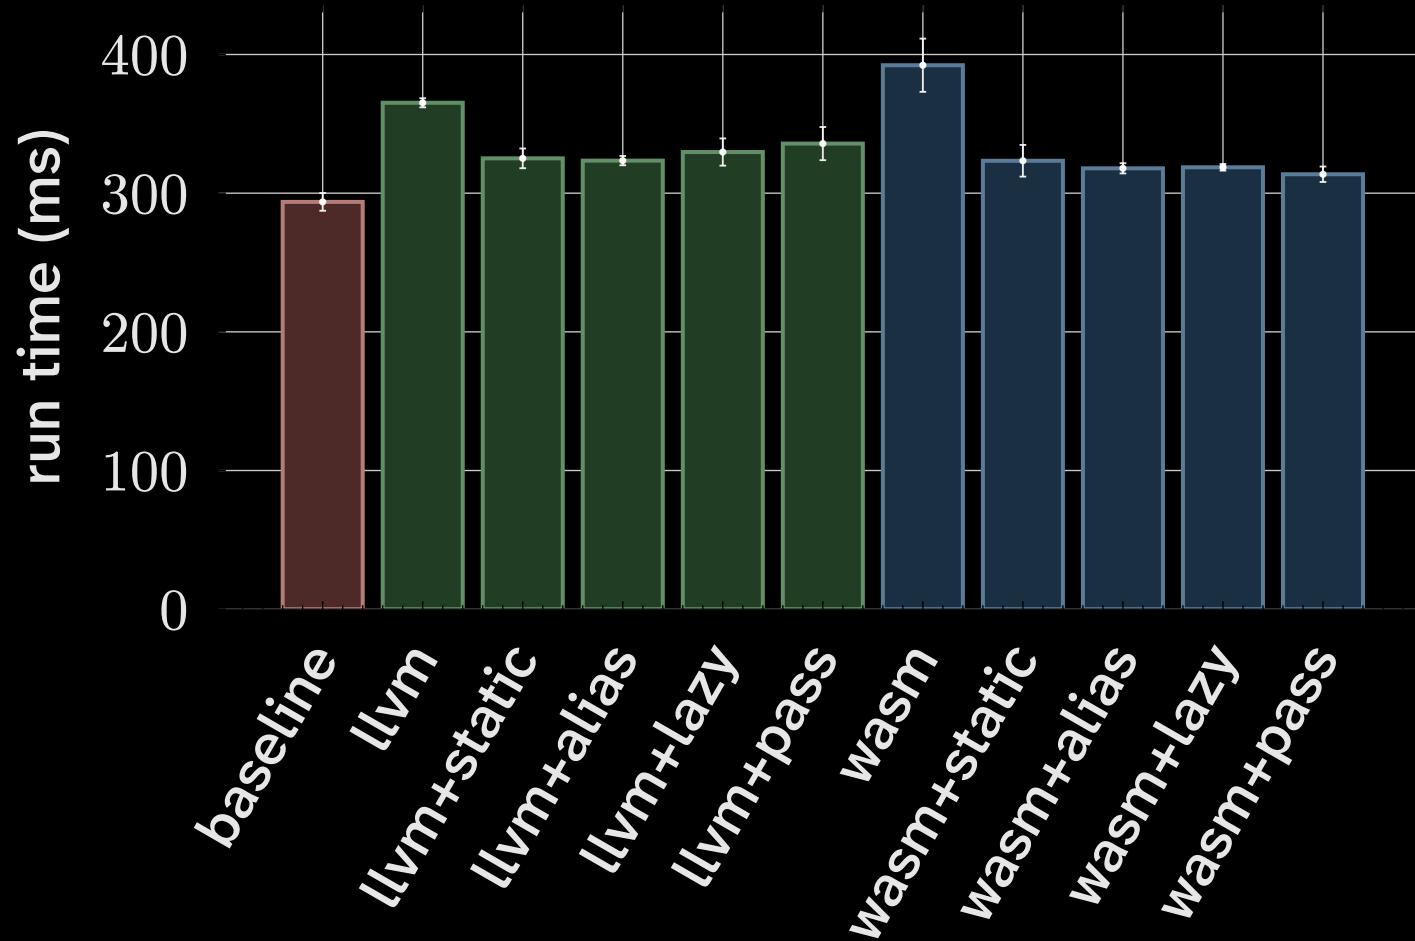
Future work: statepoints without shadow stack for native x86

Extra Slides

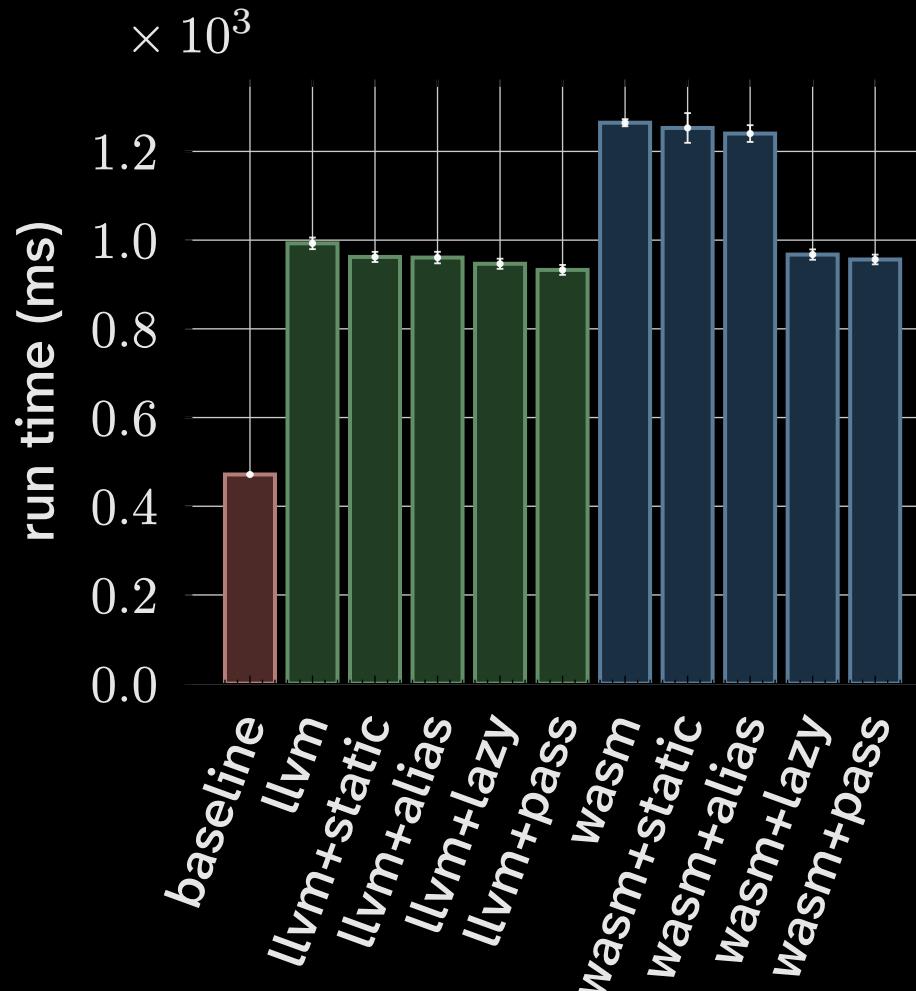
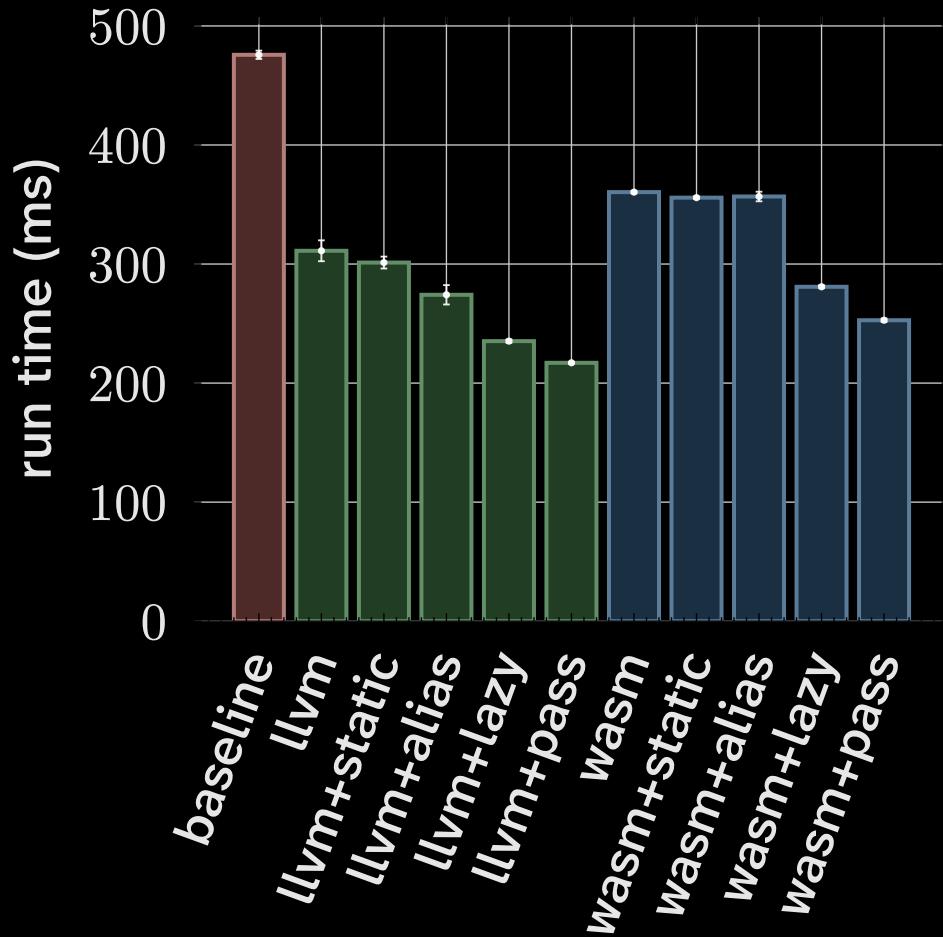
Optimisation: Static Nodes



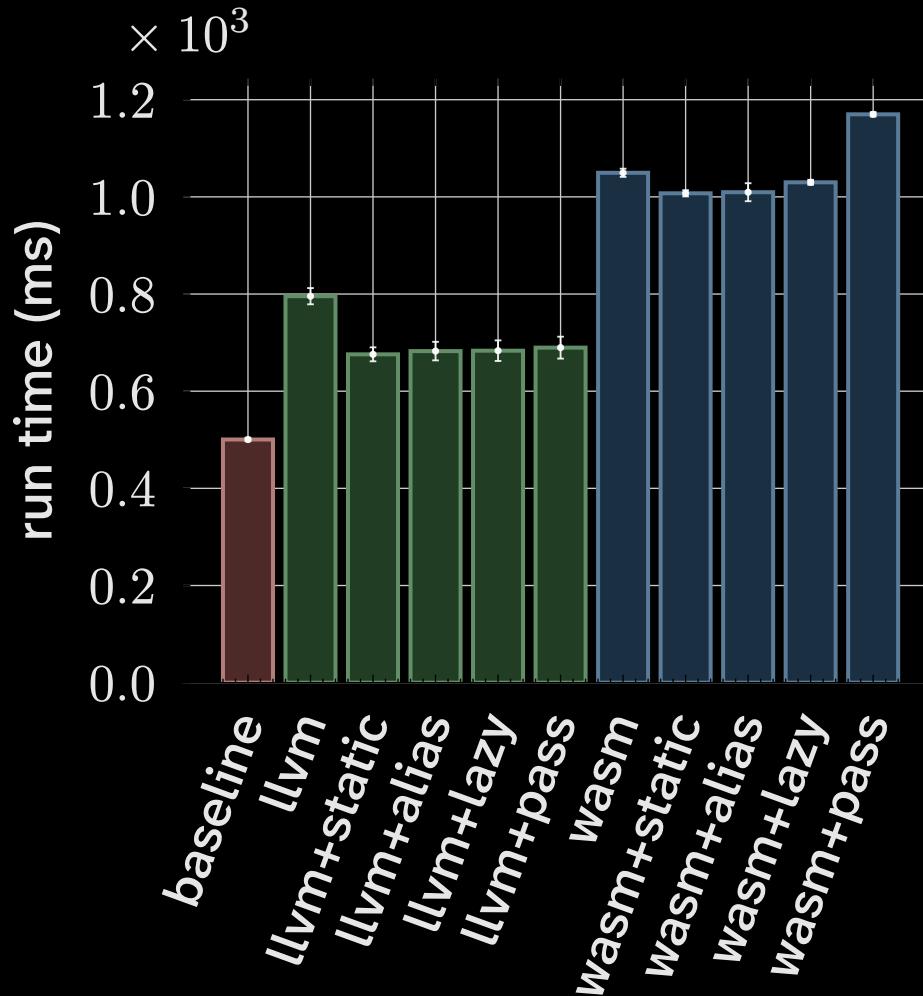
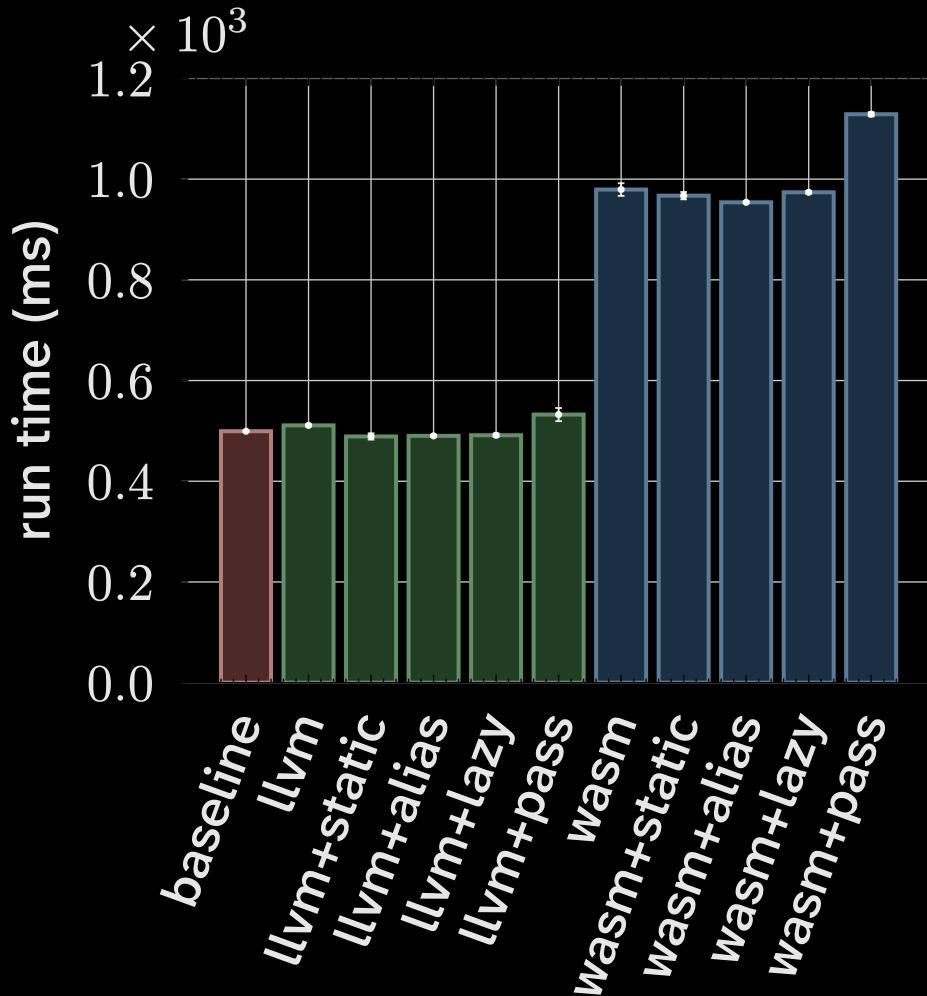
More Benchmarks: Binarytrees



More Benchmarks: Astack



More Benchmarks: List



Questions

Why not Wasm GC?

- Not supported by LLVM (yet)
- Issues with Wasm's type safety when overwriting a thunk with its result